

py-typedlogic

None

None

None

Table of contents

1. py-typedlogic: Bridging Formal Logic and Typed Python	5
1.1 Key Features	5
1.2 Installation	6
1.3 Define predicates using Pythonic idioms	6
1.4 Specify logical axioms directly in Python	6
1.5 Performing reasoning from within Python	6
2. Gallery	8
2.1 Examples	8
2.2 Synthetic Biology Design Checks with TypedLogic	10
2.3 SQL Goal Compiler	19
2.4 TLog CLI Examples	21
2.5 TLog CLI Workflow	24
3. Learning	26
3.1 Semantics: open vs closed world, NAF, WFS, and ASP	26
3.2 Multiple models	34
4. Concepts	36
4.1 typed-logic: Bridging Formal Logic and Typed Python	36
4.2 Core Concepts	36
4.3 Data Model	39
4.4 Decorators	0
4.5 Generators	0
4.6 Solvers	0
4.7 Models	0
5. Solvers	0
5.1 Solvers	0
5.2 Clingo Integrations	0
5.3 Z3 Integrations	0
5.4 Souffle Integrations	0
5.5 Prover9 Integrations	0
5.6 SnakeLog Integrations	0
5.7 ProbLog Integrations	0
5.8 Well-Founded Semantics Solver	0
5.9 Scaling the native well-founded solver	0
5.10 LLM Integrations	0

6. Compilers	0
6.1 Compilers	0
6.2 YAML Compiler	0
6.3 TLog Compiler	0
6.4 Prolog Compiler	0
6.5 ProbLog Compiler	0
6.6 TPTP Compiler	0
6.7 CLIF Compiler	0
6.8 FOR Compiler	0
6.9 SQL Compiler	0
6.10 S-Expression Compiler	0
7. Parsers	0
7.1 Parsers	0
7.2 Python Parser	0
7.3 TLog Parser	0
7.4 CLIF Parser	0
7.5 YAML Parser	0
7.6 DataFrame Parser	0
7.7 Catalog Parser	0
7.8 OWLPython Parser	0
7.9 RDF Parser	0
8. Frameworks	0
8.1 OWL-DL	0
8.2 RDF Integration	0
9. CLI Reference	0
10. typedlogic	0
10.1 convert	0
10.2 Example:	0
10.3 dump	0
10.4 Examples	0
10.5 list-compilers	0
10.6 list-parsers	0
10.7 list-solvers	0
10.8 prove	0
10.9 solve	0
10.10 Examples	0
10.11 test	0

11. Troubleshooting	0
11.1 Type Checking Errors	0
11.2 Solver Not Finding Solutions	0
11.3 Performance Issues	0
11.4 Integration Issues	0
11.5 Unexpected Logical Results	0
12. Roadmap	0
12.1 Solvers	0
12.2 Framework Integrations	0
12.3 Transformations	0
12.4 Documentation	0
12.5 Current Limitations	0

1. py-typedlogic: Bridging Formal Logic and Typed Python

TypedLogic is a powerful Python package that bridges the gap between formal logic and strongly typed Python code. It allows you to leverage fast logic programming engines like Souffle while specifying your logic in mypy-validated Python code.

links.py run.py stdout Semantics

```
# links.py
from pydantic import BaseModel
from typedlogic import FactMixin, gen2
from typedlogic.decorators import axiom

ID = str

class Link(BaseModel, FactMixin):
    """A link between two entities"""
    source: ID
    target: ID

class Path(BaseModel, FactMixin):
    """An N-hop path between two entities"""
    source: ID
    target: ID
    hops: int

@axiom
def path_from_link(x: ID, y: ID):
    """If there is a link from x to y, there is a path from x to y"""
    assert Link(source=x, target=y) >> Path(source=x, target=y, hops=1)

@axiom
def transitivity(x: ID, y: ID, z: ID, d1: int, d2: int):
    """Transitivity of paths, plus hop counting"""
    assert ((Path(source=x, target=y, hops=d1) & Path(source=y, target=z, hops=d2)) >>
            Path(source=x, target=z, hops=d1+d2))

@axiom
def reflexivity():
    """No paths back to self"""
    assert not any(Path(source=x, target=x, hops=d) for x, d in gen2(ID, int))

from typedlogic.integrations.souffle_solver import SouffleSolver
from links import Link
import links as links

solver = SouffleSolver()
solver.load(links) ## source for definitions and axioms
# Add data
links = [Link(source='CA', target='OR'), Link(source='OR', target='WA')]
for link in links:
    solver.add(link)
model = solver.model()
for fact in model.iter_retrieve("Path"):
    print(fact)
```

Output:

```
Path(source='CA', target='OR', hops=1)
Path(source='OR', target='WA', hops=1)
Path(source='CA', target='WA', hops=2)
```

To convert the Python code to first-order logic, use the CLI:

```
typedlogic convert links.py -t fol
```

Output:

```
∀[x:ID y:ID]. Link(x, y) → Path(x, y, 1)
∀[x:ID y:ID z:ID d1:int d2:int]. Path(x, y, d1) ∧ Path(y, z, d2) → Path(x, z, d1+d2)
¬∃[x:ID d:int]. Path(x, x, d)
```

1.1 Key Features

- Write logical axioms and rules using familiar Python syntax

- Benefit from strong typing and mypy validation
- Integration with multiple solvers and logic engines, including Z3 and Souffle
- Compatible with popular Python validation libraries like Pydantic

1.2 Installation

Install TypedLogic using pip:

```
pip install typedlogic
```

With all extras pre-installed:

```
pip install typedlogic[all]
```

You can also use pipx to run the [CLI](#) without installing the package globally:

```
pipx run typedlogic --help
```

1.3 Define predicates using Pythonic idioms

Inherit from one of the TypedLogic base models to add semantics to your data model. For Pydantic:

```
from typedlogic.integrations.frameworks.pydantic import FactBaseModel

ID = str

class Link(FactBaseModel):
    source: ID
    target: ID

class Path(FactBaseModel):
    source: ID
    target: ID
```

These can be used in the standard way in Python:

```
links = [Link(source='CA', target='OR'), Link(source='OR', target='WA')]
```

1.4 Specify logical axioms directly in Python

Once you have defined your data model predicates, you can specify logical axioms directly in Python:

```
from typedlogic.decorators import axiom

@axiom
def link_implies_path(x: ID, y: ID):
    """For all x, y, if there is a link from x to y, then there is a path from x to y"""
    if Link(source=x, target=y):
        assert Path(source=x, target=y)

@axiom
def transitivity(x: ID, y: ID, z: ID):
    """For all x, y, z, if there is a path from x to y and a path from y to z,
    then there is a path from x to z"""
    if Path(source=x, target=y) and Path(source=y, target=z):
        assert Path(source=x, target=z)
```

1.5 Performing reasoning from within Python

Use any of the existing solvers to perform reasoning:

```
from typedlogic.integrations.snakelogic import SnakeSolver

solver = SnakeSolver()
solver.load("links.py") ## source for definitions and axioms
for link in links:
    solver.add(link)
```

```
model = solver.model()
for fact in model.iter_retrieve("Path"):
    print(fact)
```

outputs:

```
Path(source='CA', target='OR')
Path(source='OR', target='WA')
Path(source='CA', target='WA')
```

2. Gallery

2.1 Examples

This page provides various examples to illustrate the usage of TypedLogic in different scenarios.

2.1.1 Gallery Tutorials

- [Synthetic biology design checks](#): assembly structure, GO-driven pathway feasibility, and GO-CAM curation consistency in one reusable theory.
- [SQL goal compiler](#): bind predicates to SQL tables, compile goals to SQL, run recursive CTEs, and check constraints in SQLite.
- [TLog CLI examples](#): author `.tlog` and literate `.tlog.md` files, convert them, run Clingo, filter materialized predicates, and inspect multiple worlds.
- [TLog notebook](#): notebook form of the same CLI-oriented workflow.

2.1.2 Basic Family Relationships

```
from typedlogic import FactMixin, axiom, gen2
from pydantic import BaseModel

class Parent(BaseModel, FactMixin):
    parent: str
    child: str

class Ancestor(BaseModel, FactMixin):
    ancestor: str
    descendant: str

@axiom
def parent_is_ancestor():
    return all(
        Parent(parent=x, child=y) >> Ancestor(ancestor=x, descendant=y)
        for x, y in gen2(str, str)
    )

@axiom
def ancestor_transitivity():
    return all(
        (Ancestor(ancestor=x, descendant=y) & Ancestor(ancestor=y, descendant=z))
        >> Ancestor(ancestor=x, descendant=z)
        for x, y, z in gen3(str, str, str)
    )

# Usage
theory = Theory(
    name="family_relationships",
    predicate_definitions=[
        PredicateDefinition("Parent", {"parent": str, "child": str}),
        PredicateDefinition("Ancestor", {"ancestor": str, "descendant": str}),
    ],
    sentence_groups=[SentenceGroup(name="axioms", sentences=[parent_is_ancestor(), ancestor_transitivity()])],
)

solver = Z3Solver()
solver.add(theory)
solver.add_fact(Parent(parent="Alice", child="Bob"))
solver.add_fact(Parent(parent="Bob", child="Charlie"))

result = solver.prove(Ancestor(ancestor="Alice", descendant="Charlie"))
print(f"Alice is an ancestor of Charlie: {result}")
```

2.1.3 Type Hierarchies

```
from typedlogic import FactMixin, axiom, gen1
from pydantic import BaseModel

class Animal(BaseModel, FactMixin):
    name: str

class Mammal(Animal):
    pass

class Dog(Mammal):
```



```

    breed: str

@axiom
def dogs_are_mammals():
    return all(
        Dog(name=x, breed=y) >> Mammal(name=x)
        for x, y in gen2(str, str)
    )

@axiom
def mammals_are_animals():
    return all(
        Mammal(name=x) >> Animal(name=x)
        for x in gen1(str)
    )

# Usage
theory = Theory(
    name="animal_hierarchy",
    predicate_definitions=[
        PredicateDefinition("Animal", {"name": str}),
        PredicateDefinition("Mammal", {"name": str}),
        PredicateDefinition("Dog", {"name": str, "breed": str}),
    ],
    sentence_groups=[SentenceGroup(name="axioms", sentences=[dogs_are_mammals(), mammals_are_animals()])],
)

solver = Z3Solver()
solver.add(theory)
solver.add_fact(Dog(name="Buddy", breed="Labrador"))

result = solver.prove(Animal(name="Buddy"))
print(f"Buddy is an animal: {result}")

```

These examples demonstrate how to use TypedLogic to model relationships and hierarchies, define axioms, and perform logical reasoning. You can expand on these examples to create more complex logical systems tailored to your specific use cases.

2.2 Synthetic Biology Design Checks with TypedLogic

This gallery tutorial turns a small synthetic biology design problem into composable logical checks. The same TypedLogic theory covers assembly structure, GO-driven pathway feasibility, and GO-CAM curation consistency.

The running example is sucrose utilization in *E. coli*: a design needs a transporter (`cscB`) to bring sucrose into the cell and an invertase (`cscA`) to convert intracellular sucrose into glucose and fructose.

2.2.1 Installation

This notebook uses the Clingo solver because the pathway rule is recursive and naturally evaluated as a Datalog-style fixpoint.

```
pip install 'typedlogic[clingo]'
```

2.2.2 Setup

The formal predicates and axioms live in `synbio_theory.py` next to this notebook. Keeping the theory in a normal Python module makes it reusable from tests, scripts, and notebooks.

```
from __future__ import annotations

from contextlib import contextmanager
from pathlib import Path
from typing import Iterable, Sequence
import os
import sys

from IPython.display import Markdown, display
from IPython.lib.display import Code
from typedlogic.integrations.solvers.clingo import ClingoSolver

EXAMPLE_DIR = Path.cwd()
if not (EXAMPLE_DIR / "synbio_theory.py").exists():
    EXAMPLE_DIR = Path.cwd() / "docs" / "examples"

THEORY_PATH = EXAMPLE_DIR / "synbio_theory.py"
sys.path.insert(0, str(EXAMPLE_DIR))

from synbio_theory import (
    AvailableMetabolite,
    CausalEdge,
    CausalRelation,
    DesignContains,
    EncodesProtein,
    FunctionCatalyzes,
    GOAspect,
    GOCAMIndividual,
    HasMolecularFunction,
    Part,
    RequiredMetabolite,
)

@contextmanager
def silence_clingo_stderr():
    """Suppress harmless Clingo messages about predicates that only appear as facts."""
    saved = os.dup(2)
    devnull = os.open(os.devnull, os.O_WRONLY)
    os.dup2(devnull, 2)
    try:
        yield
    finally:
        os.dup2(saved, 2)
        os.close(saved)
        os.close(devnull)

def infer(facts: Iterable[object], *predicate_names: str) -> dict[str, list[tuple[object, ...]]]:
    """Load the synbio theory, add facts, and return selected derived predicates."""
    solver = ClingoSolver()
    solver.load(THEORY_PATH)
    for fact in facts:
        solver.add_fact(fact)

    with silence_clingo_stderr():
        model = solver.model()

    rows = {predicate_name: [] for predicate_name in predicate_names}
    for term in model.ground_terms:
```

```

        if term.predicate in rows:
            rows[term.predicate].append(term.values)
        return {predicate: sorted(values) for predicate, values in rows.items()}

def markdown_table(headers: Sequence[str], rows: Sequence[Sequence[object]]) -> str:
    """Render a small Markdown table without adding a pandas dependency."""
    lines = ["| " + " | ".join(headers) + " |"]
    lines.append("| " + " | ".join(["---"] * len(headers)) + " |")
    for row in rows:
        lines.append("| " + " | ".join(str(value) for value in row) + " |")
    return "\n".join(lines)

def pairs(values: Sequence[tuple[object, ...]]) -> str:
    """Format binary or assembly-scoped predicate values as compact edges."""
    formatted = []
    for value in values:
        if len(value) == 2:
            left, right = value
            formatted.append(f"{left}->{right}")
        elif len(value) == 3:
            assembly, left, right = value
            formatted.append(f"{assembly}:{left}->{right}")
        else:
            formatted.append("->".join(str(part) for part in value))
    return ", ".join(formatted) or "-"

def csv(values: Sequence[object]) -> str:
    """Format a short sequence for a Markdown table cell."""
    return ", ".join(str(value) for value in values) or "-"

```

```
Code(filename=str(THEORY_PATH), language="python")
```

```

"""Synthetic biology design checks for the gallery notebook."""

# ruff: noqa: S101
#
# The theory is intentionally small, but it spans three practical layers:
# 1. Assembly structure checks for compatible Golden Gate-style overhangs.
# 2. GO-driven pathway reachability for engineered metabolic functions.
# 3. GO-CAM curation checks for causal edges between molecular functions.

from dataclasses import dataclass

from typedlogic import Fact, axiom

@dataclass(frozen=True)
class Part(Fact):
    """A DNA part in an ordered assembly."""

    assembly: str
    name: str
    part_type: str
    five_oh: str
    three_oh: str
    position: int

@dataclass(frozen=True)
class CanLigate(Fact):
    """Two parts have compatible overhangs and can ligate."""

    assembly: str
    upstream: str
    downstream: str

@dataclass(frozen=True)
class IntendedAdjacent(Fact):
    """Two parts are adjacent in the intended design order."""

    assembly: str
    upstream: str
    downstream: str

@dataclass(frozen=True)
class IntendedLigationGap(Fact):
    """Two intended adjacent parts have incompatible overhangs."""

```

```

assembly: str
upstream: str
downstream: str

@dataclass(frozen=True)
class MisligationRisk(Fact):
    """A compatible ligation exists outside the intended adjacency order."""

    assembly: str
    upstream: str
    downstream: str

@axiom
def ligation_compatibility(
    assembly1: str,
    n1: str,
    t1: str,
    fo1: str,
    to1: str,
    pos1: int,
    assembly2: str,
    n2: str,
    t2: str,
    fo2: str,
    to2: str,
    pos2: int,
):
    """Derive ligation compatibility when the upstream 3' overhang matches the downstream 5' overhang."""
    if Part(assembly1, n1, t1, fo1, to1, pos1) and Part(assembly2, n2, t2, fo2, to2, pos2):
        if assembly1 == assembly2 and to1 == fo2 and n1 != n2:
            assert CanLigate(assembly1, n1, n2)

@axiom
def intended_adjacency(
    assembly1: str,
    n1: str,
    t1: str,
    fo1: str,
    to1: str,
    pos1: int,
    assembly2: str,
    n2: str,
    t2: str,
    fo2: str,
    to2: str,
    pos2: int,
):
    """Derive intended adjacency from the declared part positions."""
    if Part(assembly1, n1, t1, fo1, to1, pos1) and Part(assembly2, n2, t2, fo2, to2, pos2):
        if assembly1 == assembly2 and pos2 == pos1 + 1:
            assert IntendedAdjacent(assembly1, n1, n2)

@axiom
def intended_ligation_gap(
    assembly1: str,
    n1: str,
    t1: str,
    fo1: str,
    to1: str,
    pos1: int,
    assembly2: str,
    n2: str,
    t2: str,
    fo2: str,
    to2: str,
    pos2: int,
):
    """Flag intended adjacent parts whose overhangs do not match."""
    if Part(assembly1, n1, t1, fo1, to1, pos1) and Part(assembly2, n2, t2, fo2, to2, pos2):
        if assembly1 == assembly2 and pos2 == pos1 + 1 and to1 != fo2:
            assert IntendedLigationGap(assembly1, n1, n2)

@axiom
def misligation_detection(
    assembly: str,

```

```

upstream: str,
downstream: str,
n1: str,
t1: str,
fo1: str,
to1: str,
pos1: int,
n2: str,
t2: str,
fo2: str,
to2: str,
pos2: int,
):
    """Flag compatible ligations that skip over the intended next position."""
    if (
        CanLigate(assembly, upstream, downstream)
        and Part(assembly, n1, t1, fo1, to1, pos1)
        and Part(assembly, n2, t2, fo2, to2, pos2)
        and upstream == n1
        and downstream == n2
    ):
        if pos2 != pos1 + 1:
            assert MisligationRisk(assembly, upstream, downstream)

@dataclass(frozen=True)
class EncodesProtein(Fact):
    """A CDS part encodes a protein."""

    part: str
    protein: str

@dataclass(frozen=True)
class HasMolecularFunction(Fact):
    """A protein has a GO molecular function."""

    protein: str
    go_mf: str

@dataclass(frozen=True)
class FunctionCatalyzes(Fact):
    """A molecular function converts a substrate metabolite to a product metabolite."""

    go_mf: str
    substrate: str
    product: str

@dataclass(frozen=True)
class DesignContains(Fact):
    """A named design contains a part."""

    design: str
    part: str

@dataclass(frozen=True)
class AvailableMetabolite(Fact):
    """A metabolite is available to a design from the environment or an encoded function."""

    design: str
    metabolite: str

@dataclass(frozen=True)
class RequiredMetabolite(Fact):
    """A metabolite that the engineered cell needs for the design intent."""

    design: str
    metabolite: str

@axiom
def metabolite_produced_by_design(
    design: str,
    part: str,
    protein: str,
    go_mf: str,

```

```

    substrate: str,
    product: str,
):
    """Propagate metabolite availability through functions encoded by the design."""
    if (
        DesignContains(design, part)
        and EncodesProtein(part, protein)
        and HasMolecularFunction(protein, go_mf)
        and FunctionCatalyzes(go_mf, substrate, product)
        and AvailableMetabolite(design, substrate)
    ):
        assert AvailableMetabolite(design, product)

@dataclass(frozen=True)
class GOAspect(Fact):
    """The aspect of a GO term: molecular function, biological process, or cellular component."""

    go_term: str
    aspect: str

@dataclass(frozen=True)
class GOCAMIndividual(Fact):
    """An individual in a GO-CAM model instantiating a GO class."""

    iri: str
    go_class: str

@dataclass(frozen=True)
class CausalRelation(Fact):
    """A relation that should connect molecular function individuals in GO-CAM."""

    relation: str

@dataclass(frozen=True)
class CausalEdge(Fact):
    """A causal relation between two GO-CAM individuals."""

    upstream_iri: str
    downstream_iri: str
    relation: str

@dataclass(frozen=True)
class GOCAMViolation(Fact):
    """A derived GO-CAM curation rule violation."""

    individual: str
    rule: str

@axiom
def causal_edge_upstream_must_be_mf(
    up: str,
    down: str,
    rel: str,
    upstream_class: str,
    aspect: str,
):
    """Require the upstream node of a causal edge to instantiate a molecular function."""
    if (
        CausalEdge(up, down, rel)
        and CausalRelation(rel)
        and GOCAMIndividual(up, upstream_class)
        and GOAspect(upstream_class, aspect)
    ):
        if aspect != "MF":
            assert GOCAMViolation(up, "causal_upstream_not_MF")

@axiom
def causal_edge_downstream_must_be_mf(
    up: str,
    down: str,
    rel: str,
    downstream_class: str,
    aspect: str,

```

```

):
    """Require the downstream node of a causal edge to instantiate a molecular function."""
    if (
        CausalEdge(up, down, rel)
        and CausalRelation(rel)
        and GOCAMIndividual(down, downstream_class)
        and GOAspect(downstream_class, aspect)
    ):
        if aspect != "MF":
            assert GOCAMViolation(down, "causal_downstream_not_MF")

```

2.2.3 Layer 1: Assembly Structure

The structural layer treats overhangs as assembly-scoped typed facts. From those facts it derives which parts can ligate, which parts are intended neighbors, whether an intended neighboring pair fails to ligate, and whether any compatible ligation skips the intended next position.

```

safe_assembly = [
    Part("clean", "p_lac", "promoter", "GGAG", "TACT", 0),
    Part("clean", "rbs_b0034", "rbs", "TACT", "AATG", 1),
    Part("clean", "cscB", "cds", "AATG", "GCTT", 2),
    Part("clean", "term_b0015", "terminator", "GCTT", "CGCT", 3),
]

risky_assembly = [
    Part("risky", "p_lac", "promoter", "GGAG", "TACT", 0),
    Part("risky", "rbs_b0034", "rbs", "TACT", "AATG", 1),
    Part("risky", "cscB_skip_rbs", "cds", "TACT", "GCTT", 2),
    Part("risky", "term_b0015", "terminator", "GCTT", "CGCT", 3),
]

def assembly_report(label: str, facts: Sequence[Part]) -> tuple[str, str, str, str, str]:
    results = infer(facts, "CanLigate", "IntendedAdjacent", "IntendedLigationGap", "MisligationRisk")
    gaps = results["IntendedLigationGap"]
    risks = results["MisligationRisk"]
    return (
        label,
        pairs(results["CanLigate"]),
        pairs(results["IntendedAdjacent"]),
        pairs(gaps) if gaps else "none",
        pairs(risks) if risks else "none",
    )

assembly_rows = [
    assembly_report("clean overhang grammar", safe_assembly),
    assembly_report("CDS can skip RBS", risky_assembly),
]

display(
    Markdown(
        markdown_table(
            ["Assembly", "Can ligate", "Intended adjacency", "Intended gap", "Skip risk"],
            assembly_rows,
        )
    )
)

```

Assembly	Can ligate	Intended adjacency	Intended gap	Skip risk
clean overhang grammar	clean:cscB->term_b0015, clean:p_lac->rbs_b0034, clean:rbs_b0034->cscB	clean:cscB->term_b0015, clean:p_lac->rbs_b0034, clean:rbs_b0034->cscB	none	none
CDS can skip RBS	risky:cscB_skip_rbs->term_b0015, risky:p_lac->cscB_skip_rbs, risky:p_lac->rbs_b0034	risky:cscB_skip_rbs->term_b0015, risky:p_lac->rbs_b0034, risky:rbs_b0034->cscB_skip_rbs	risky:rbs_b0034->cscB_skip_rbs	risky:p_lac->cscB_skip_rbs

In the second assembly, `rbs_b0034` and `cscB_skip_rbs` are intended neighbors but have incompatible overhangs. The same assembly also lets `p_lac` ligate directly into `cscB_skip_rbs`, so the theory reports both the intended ligation gap and the skip risk before considering pathway function.

2.2.4 Layer 2: GO-Driven Pathway Feasibility

The functional layer connects parts to proteins, proteins to GO molecular functions, and GO functions to metabolite transformations. A recursive rule derives every metabolite reachable from the growth medium.

```
BIOLOGY = [
    EncodesProtein("cscA", "P10000_cscA"),
    EncodesProtein("cscB", "P30000_cscB"),
    HasMolecularFunction("P10000_cscA", "GO:0004575"), # sucrose alpha-glucosidase
    HasMolecularFunction("P30000_cscB", "GO:0008515"), # sucrose:H+ symporter
    FunctionCatalyzes("GO:0008515", "sucrose_out", "sucrose_in"),
    FunctionCatalyzes("GO:0004575", "sucrose_in", "glucose_in"),
    FunctionCatalyzes("GO:0004575", "sucrose_in", "fructose_in"),
]

def pathway_report(label: str, design_parts: Sequence[str]) -> tuple[str, str, str, str]:
    facts = [
        *BIOLOGY,
        *(DesignContains("sucrose_design", part) for part in design_parts),
        AvailableMetabolite("sucrose_design", "sucrose_out"),
        RequiredMetabolite("sucrose_design", "glucose_in"),
    ]
    results = infer(facts, "AvailableMetabolite", "RequiredMetabolite")
    available = sorted({metabolite for _, metabolite in results["AvailableMetabolite"]})
    required = sorted({metabolite for _, metabolite in results["RequiredMetabolite"]})
    missing = [metabolite for metabolite in required if metabolite not in available]
    status = "complete" if not missing else f"gap: {csv(missing)}"
    return label, csv(design_parts), csv(available), status

pathway_rows = [
    pathway_report("full sucrose pathway", ["cscA", "cscB"]),
    pathway_report("permease only", ["cscB"]),
    pathway_report("invertase only", ["cscA"]),
    pathway_report("empty design", []),
]

display(Markdown(markdown_table(["Design", "Parts", "Reachable metabolites", "Verdict"], pathway_rows)))
```

Design	Parts	Reachable metabolites	Verdict
full sucrose pathway	cscA, cscB	fructose_in, glucose_in, sucrose_in, sucrose_out	complete
permease only	cscB	sucrose_in, sucrose_out	gap: glucose_in
invertase only	cscA	sucrose_out	gap: glucose_in
empty design	-	sucrose_out	gap: glucose_in

The full design reaches intracellular glucose from extracellular sucrose. The permease-only design imports sucrose but cannot cleave it. The invertase-only design has the enzyme, but no route for sucrose to enter the cell.

2.2.5 Layer 3: GO-CAM Curation Consistency

The curation layer checks a common GO-CAM modeling rule: declared causal relations such as `R0:0002413` should connect molecular function individuals, not biological process individuals.

```
GO_ASPECTS = [
    GOAspect("GO:0004575", "MF"), # sucrose alpha-glucosidase activity
    GOAspect("GO:0008515", "MF"), # sucrose:H+ symporter activity
    GOAspect("GO:0005992", "BP"), # trehalose biosynthetic process
    GOAspect("GO:0006012", "BP"), # galactose metabolic process
]

CAUSAL_RELATIONS = [
    CausalRelation("R0:0002411"), # causally upstream of
    CausalRelation("R0:0002413"), # provides direct input for
]

def gocam_report(
    label: str,
    individuals: Sequence[tuple[str, str]],
    edges: Sequence[tuple[str, str, str]],
) -> tuple[str, str, str, str]:
    facts = [
        *GO_ASPECTS,
        *CAUSAL_RELATIONS,
        *(GOCAMIndividual(iri, go_class) for iri, go_class in individuals),
        *(CausalEdge(upstream, downstream, relation) for upstream, downstream, relation in edges),
    ]
```



```

}
violations = infer(facts, "GOCAMViolation")["GOCAMViolation"]
violation_text = ", ".join(f"{iri}: {rule}" for iri, rule in violations) or "none"
return (
    label,
    csv([f"{iri}={go_class}" for iri, go_class in individuals]),
    pairs([(upstream, downstream) for upstream, downstream, _ in edges]),
    violation_text,
)

```

```

gocam_rows = [
    gocam_report(
        "valid MF to MF causal edge",
        individuals=[("ind:1", "GO:0008515"), ("ind:2", "GO:0004575")],
        edges=[("ind:1", "ind:2", "R0:0002413")],
    ),
    gocam_report(
        "invalid MF to BP causal edge",
        individuals=[("ind:3", "GO:0008515"), ("ind:4", "GO:0005992")],
        edges=[("ind:3", "ind:4", "R0:0002413")],
    ),
    gocam_report(
        "invalid BP to BP causal edge",
        individuals=[("ind:5", "GO:0006012"), ("ind:6", "GO:0005992")],
        edges=[("ind:5", "ind:6", "R0:0002413")],
    ),
    gocam_report(
        "non-causal relation to BP",
        individuals=[("ind:7", "GO:0008515"), ("ind:8", "GO:0005992")],
        edges=[("ind:7", "ind:8", "BF0:0000050")],
    ),
]

display(Markdown(markdown_table(["GO-CAM model", "Individuals", "Causal edge", "Violations"], gocam_rows)))

```

GO-CAM model	Individuals	Causal edge	Violations
valid MF to MF causal edge	ind:1=GO:0008515, ind:2=GO:0004575	ind:1->ind:2	none
invalid MF to BP causal edge	ind:3=GO:0008515, ind:4=GO:0005992	ind:3->ind:4	ind:4: causal_downstream_not_MF
invalid BP to BP causal edge	ind:5=GO:0006012, ind:6=GO:0005992	ind:5->ind:6	ind:5: causal_upstream_not_MF, ind:6: causal_downstream_not_MF
non-causal relation to BP	ind:7=GO:0008515, ind:8=GO:0005992	ind:7->ind:8	none

The same solver workflow handles design-time checks and curation-time checks. In a production pipeline, the facts could come from a parts registry, GO annotations, Rhea mappings, and GO-CAM RDF exports rather than hand-authored lists.

2.2.6 Inspect the Compiled Rules

TypedLogic parses the Python axioms into a logical theory and the Clingo backend renders that theory as ASP. The recursive pathway rule is the key fixpoint rule.

```

solver = ClingoSolver()
solver.load(THEORY_PATH)
compiled_rules = solver.dump().splitlines()
for rule in compiled_rules:
    if rule.startswith(("availablemetabolite", "intendedligationgap", "misligationrisk", "gocamviolation")):
        print(rule)

```

```

intendedligationgap(Assembly1, N1, N2) :- part(Assembly1, N1, T1, Fo1, To1, Pos1), part(Assembly2, N2, T2, Fo2, To2, Pos2), Assembly1 == Assembly2, Pos2 == Pos1 + 1, To1 != Fo2.
misligationrisk(Assembly, Upstream, Downstream) :- canligate(Assembly, Upstream, Downstream), part(Assembly, N1, T1, Fo1, To1, Pos1), part(Assembly, N2, T2, Fo2, To2, Pos2), Upstream == N1, Downstream == N2, Pos2 != Pos1 + 1.
availablemetabolite(Design, Product) :- designcontains(Design, Part), encodesprotein(Part, Protein), hasmolecularfunction(Protein, Go_mf), functioncatalyzes(Go_mf, Substrate, Product), availablemetabolite(Design, Substrate).
gocamviolation(Up, "causal_upstream_not_MF") :- causededge(Up, Down, Rel), causalrelation(Rel), gocamindividual(Up, Upstream_class), goaspect(Upstream_class, Aspect), Aspect != "MF".
gocamviolation(Down, "causal_downstream_not_MF") :- causededge(Up, Down, Rel), causalrelation(Rel), gocamindividual(Down, Downstream_class), goaspect(Downstream_class, Aspect), Aspect != "MF".

```

2.2.7 Next Steps

This pattern scales by swapping the toy facts for real sources: a parts library for assembly facts, GO and Rhea mappings for function-to-reaction facts, and GO-CAM exports for curation facts. Violations can then be reported back as design review comments or curator feedback.

2.3 SQL Goal Compiler

This notebook demonstrates how to translate a TypedLogic theory plus a query or goal into SQL.

The examples use an in-memory SQLite database so the generated SQL is executed end to end.

2.3.1 Setup

The SQL compiler is a regular TypedLogic compiler. `PredicateBinding` tells the compiler which SQL table and columns back an extensional predicate.

```
import sqlite3

from typedlogic import And, Forall, Implies, NegationAsFailure, Not, PredicateDefinition, SentenceGroup, Term, Theory, Variable
from typedlogic.compilers.sql_compiler import PredicateBinding, SQLCompiler, SQLCompilerConfig
from typedlogic.datamodel import SentenceGroupType

def fetchall(sql: str, setup: list[str] | None = None) -> list[tuple]:
    """Run SQL in a fresh in-memory SQLite database."""
    connection = sqlite3.connect(":memory:")
    try:
        for statement in setup or []:
            connection.execute(statement)
        return list(connection.execute(sql))
    finally:
        connection.close()

compiler = SQLCompiler()
```

2.3.2 Rules As Views

This theory defines `Ancestor` from `Parent`. The recursive rule is compiled as a recursive common table expression.

```
x = Variable("x")
y = Variable("y")
z = Variable("z")

family_theory = Theory(
    predicate_definitions=[
        PredicateDefinition("Parent", {"parent": "str", "child": "str"}),
        PredicateDefinition("Ancestor", {"ancestor": "str", "descendant": "str"}),
    ]
)
family_theory.add(Forall([x, y], Implies(Term("Parent", x, y), Term("Ancestor", x, y))))
family_theory.add(
    Forall(
        [x, y, z],
        Implies(And(Term("Ancestor", x, y), Term("Parent", y, z)), Term("Ancestor", x, z)),
    )
)

ancestor_sql = compiler.compile(
    family_theory,
    goal=Term("Ancestor", "Alice", Variable("who")),
    bindings={"Parent": PredicateBinding("parent_edges", {"parent": "src", "child": "dst"})},
)
print(ancestor_sql)
```

```
fetchall(
    ancestor_sql,
    [
        "CREATE TABLE parent_edges (src TEXT, dst TEXT)",
        "INSERT INTO parent_edges VALUES ('Alice', 'Bob'), ('Bob', 'Carol'), ('Carol', 'Dana')",
    ],
)
```

2.3.3 Inline Unit Clauses

Ground facts do not require table bindings. They become inline CTE branches.

```
person_theory = Theory(predicate_definitions=[PredicateDefinition("Person", {"name": "str"})])
person_theory.add(Term("Person", "Alice"))
person_theory.add(Term("Person", "Bob"))
```

```

person_sql = compiler.compile(person_theory, goal=Term("Person", Variable("name")))
print(person_sql)
fetchall(person_sql)

```

2.3.4 Goal Sentence Groups

If no explicit `goal` is supplied, the compiler uses goal sentence groups from the theory.

```

goal_theory = Theory(predicate_definitions=[PredicateDefinition("Person", {"name": "str"})])
goal_theory.add(Term("Person", "Alice"))
goal_theory.sentence_groups.append(
    SentenceGroup(name="goals", group_type=SentenceGroupType.GOAL, sentences=[Term("Person", Variable("name"))])
)

goal_sql = compiler.compile(goal_theory)
print(goal_sql)
fetchall(goal_sql.removesuffix(";"))

```

2.3.5 Negation As Failure

Negated body atoms are emitted as correlated `NOT EXISTS` subqueries.

```

eligibility_theory = Theory(
    predicate_definitions=[
        PredicateDefinition("Person", {"name": "str"}),
        PredicateDefinition("Blocked", {"name": "str"}),
        PredicateDefinition("Eligible", {"name": "str"}),
    ]
)
eligibility_theory.add(
    Forall([x], Implies(And(Term("Person", x), NegationAsFailure(Term("Blocked", x))), Term("Eligible", x)))
)

eligible_sql = compiler.compile(
    eligibility_theory,
    goal=Term("Eligible", Variable("name")),
    bindings={"Person": "people", "Blocked": "blocked"},
)
print(eligible_sql)
fetchall(
    eligible_sql,
    [
        "CREATE TABLE people (name TEXT)",
        "CREATE TABLE blocked (name TEXT)",
        "INSERT INTO people VALUES ('Alice'), ('Bob')",
        "INSERT INTO blocked VALUES ('Bob')",
    ],
)

```

2.3.6 Constraint Checks

First-order denial constraints compile to validation queries. A returned row is a violation witness.

```

constraint_theory = Theory(
    predicate_definitions=[
        PredicateDefinition("Cat", {"name": "str"}),
        PredicateDefinition("Dog", {"name": "str"}),
    ]
)
constraint_theory.add(Forall([x], Not(And(Term("Cat", x), Term("Dog", x)))))

[constraint_sql] = compiler.compile_constraints(
    constraint_theory,
    config=SQLCompilerConfig(bindings={"Cat": PredicateBinding("cats"), "Dog": PredicateBinding("dogs")}),
)
print(constraint_sql)
fetchall(
    constraint_sql,
    [
        "CREATE TABLE cats (name TEXT)",
        "CREATE TABLE dogs (name TEXT)",
        "INSERT INTO cats VALUES ('Jules'), ('Milo')",
        "INSERT INTO dogs VALUES ('Milo'), ('Rex')",
    ],
)

```

2.4 TLog CLI Examples

TLog is a parser and compiler format, not a separate CLI mode. Use the existing generic commands:

- `convert` for one input theory to one output format
- `dump` for combining and exporting inputs
- `solve` for satisfiability, materialization, and answer-set enumeration
- `test` for quoted `test_case(...)` metadata
- `prove` for quoted goals and lemmas

2.4.1 Convert TLog To Other Formats

```
typedlogic convert docs/examples/tlog/ancestor.tlog -t prolog
```

Example output:

```
%% Predicate Definitions
% parent(parent: PersonID, child: PersonID)
% ancestor(ancestor: PersonID, descendant: PersonID)

parent('Alice', 'Bob').
parent('Bob', 'Charlie').
ancestor(X, Y) :- parent(X, Y).
ancestor(X, Z) :- ancestor(X, Y), ancestor(Y, Z).
```

Other compiler targets work the same way:

```
typedlogic convert docs/examples/tlog/ancestor.tlog -t yaml
typedlogic convert docs/examples/tlog/ancestor.tlog -t fol
typedlogic convert docs/examples/tlog/ancestor.tlog -t souffle
```

2.4.2 Convert To TLog

Use `-t tlog` with any parser-supported input:

```
typedlogic convert docs/examples/tlog/ancestor.tlog -t tlog
```

Example output:

```
type PersonID: str.

pred parent(parent: PersonID, child: PersonID).
pred ancestor(ancestor: PersonID, descendant: PersonID).

parent('Alice', 'Bob').
parent('Bob', 'Charlie').
/// Direct parent links are ancestor links.
all x, y | ancestor(x, y) :- parent(x, y).
/// Ancestor links are transitive.
all x, y, z | ancestor(x, z) :- (ancestor(x, y) & ancestor(y, z)).
```

2.4.3 Literate Markdown

Files ending in `.tlog.md` are parsed as Markdown wrappers around fenced TLog blocks:

```
typedlogic dump docs/examples/tlog/literate-rules.tlog.md -t prolog
```

The prose is ignored. Fenced `tlog`, `typedlogic`, or `logic` blocks are parsed in order.

2.4.4 Run Inference With Clingo

```
typedlogic solve docs/examples/tlog/ancestor.tlog \
--solver clingo \
```

```
--show ancestor \
--max-models 1
```

Example output:

```
Satisfiable: True

=== Model 1 ===
ancestor(Alice, Bob)
ancestor(Bob, Charlie)
ancestor(Alice, Charlie)

Total models shown: 1
```

`--show ancestor` filters materialized output to selected predicates. It is a generic `solve` option and works for other syntaxes too.

To inspect the generated solver-specific program before solving, use `--dump-program`:

```
typedlogic solve docs/examples/tlog/ancestor.tlog \
--solver clingo \
--dump-program
```

2.4.5 Show Multiple Worlds

Answer-set solvers such as Clingo can produce multiple models:

```
typedlogic solve docs/examples/tlog/worlds.tlog \
--solver clingo \
--show selected \
--max-models 2
```

Example output:

```
Satisfiable: True

=== Model 1 ===
selected(coffee)

=== Model 2 ===
selected(tea)

Total models shown: 2
```

The input syntax is still just TLog; the multiple-world behavior comes from the chosen solver.

2.4.6 Run Quoted Tests

TLog files can include test cases as quoted metadata:

```
pred human(name: str).
pred mortal(name: str).

mortal(x) :- human(x).

test_case(
  "socrates_mortality",
  given(that(human("socrates"))),
  expect(that(satisfiable() & mortal("socrates") & not philosopher("socrates")))
).
```

`solve` ignores test cases, goals, and lemmas. Run validation explicitly with `test`; it checks test cases and also proves goals and lemmas:

```
typedlogic test docs/examples/tlog/mortality.tlog --solver clingo
typedlogic test docs/examples/tlog/mortality.tlog --solver clingo --test socrates_mortality
```

Example output:

```
PASS socrates_mortality
1 test case(s), 0 failed, 0 unknown
```

Use `--dump-program` to print the generated solver program for each test fixture before expectations are checked. Use `--no-proofs` when you only want test-case fixtures and expectations.

2.4.7 Prove Lemmas

Lemmas are quoted proof obligations, not axioms:

```
lemma("socrates_is_mortal", that(mortal("socrates"))).
```

`test` proves proof obligations by default. Use `prove` for proof-only runs:

```
typedlogic prove docs/examples/tlog/mortality.tlog --solver z3
typedlogic prove docs/examples/tlog/mortality.tlog --solver z3 --target lemmas
typedlogic prove docs/examples/tlog/mortality.tlog --solver z3 --name socrates_is_mortal
```

Example output:

```
PASS lemma socrates_is_mortal: mortal('socrates')
1 obligation(s), 0 failed, 0 unknown
```

2.5 TLog CLI Workflow

This notebook shows the command-line workflow for TLog. TLog is just another parser/compiler syntax: the generic `typedlogic convert`, `typedlogic dump`, and `typedlogic solve` commands do the work.

```
from pathlib import Path

workspace = Path("tlog-demo")
workspace.mkdir(exist_ok=True)
ancestor_tlog = workspace / "ancestor.tlog"
ancestor_tlog.write_text(
    """
type PersonID: str.

pred parent(parent: PersonID, child: PersonID).
pred ancestor(ancestor: PersonID, descendant: PersonID).

parent("Alice", "Bob").
parent("Bob", "Charlie").

/// Direct parent links are ancestor links.
ancestor(x, y) :- parent(x, y).

/// Ancestor links are transitive.
ancestor(x, z) :- ancestor(x, y), ancestor(y, z).
    """.strip()
    + "\n",
    encoding="utf-8",
)
ancestor_tlog
```

2.5.1 Convert TLog To Other Formats

The input format is inferred from `.tlog`; the output format is selected with `-t`.

```
!typedlogic convert tlog-demo/ancestor.tlog -t prolog
```

```
!typedlogic convert tlog-demo/ancestor.tlog -t yaml
```

2.5.2 Convert Back To TLog

`tlog` is also a compiler target.

```
!typedlogic convert tlog-demo/ancestor.tlog -t tlog
```

2.5.3 Run Inference With Clingo

`--show` filters materialized ground terms by predicate. `--max-models` limits model output.

```
!typedlogic solve tlog-demo/ancestor.tlog --solver clingo --show ancestor --max-models 1
```

2.5.4 Literate Markdown

A `.tlog.md` file can mix prose and fenced TLog blocks.

```
literate = workspace / "literate-rules.tlog.md"
literate.write_text(
    """
# Family Rules

Only fenced `tlog` blocks are parsed.

``tlog
pred parent(parent: str, child: str).
pred ancestor(ancestor: str, descendant: str).
parent("Alice", "Bob").
∀ x, y | parent(x, y) -> ancestor(x, y).
    """.strip() + "\n", encoding="utf-8", )
literate
```



```

</div>
</div>
<div class="cell border-box-sizing code_cell rendered" markdown="1">
<div class="input">

```python
!typedlogic dump tlog-demo/literate-rules.tlog.md -t prolog

```

## 2.5.5 Multiple Worlds

When the selected solver supports multiple models, the same generic `solve` command can show them.

```

worlds_tlog = workspace / "worlds.tlog"
worlds_tlog.write_text(
 """
pred selected(option: str).
pred hidden(option: str).

selected("tea") | selected("coffee").
hidden("noise").
 """.strip()
 + "\n",
 encoding="utf-8",
)
worlds_tlog

```

```
!typedlogic solve tlog-demo/worlds.tlog --solver clingo --show selected --max-models 2
```

## 3. Learning

### 3.1 Semantics: open vs closed world, NAF, WFS, and ASP

One innocent-looking symbol — *negation* — hides some of the deepest design choices in logic programming. When you write `"not Flies(tweety) "`, what does it mean?

- Is it **false**, or merely **not known to be true**?
- If you learn a new fact tomorrow, can a conclusion you drew today be **withdrawn**?
- If two conclusions are locked in a mutual "each is true unless the other is", is the answer **two possible worlds, no world at all**, or **"undefined"**?

Different logics answer these questions differently. This notebook is an *executable* tour of those answers. Every claim below is backed by a solver you can re-run — no hand-waving.

We author each theory in **TLog**, the compact text syntax for typedlogic, so the logic reads directly instead of hiding inside Python strings. The same theory files are loaded and handed to different solvers.

We cover four ideas and the relationships between them:

Idea	One-line intuition	Illustrated with
<b>Open-World Assumption (OWA)</b>	absence of a fact means <i>unknown</i>	Z3 (classical FOL)
<b>Closed-World Assumption (CWA)</b>	absence of a fact means <i>false</i>	Clingo, Datalog
<b>Negation as Failure (NAF)</b>	<code>not p</code> succeeds if <code>p</code> can't be proven	Clingo
<b>Answer Set Programming (ASP)</b>	stable models: enumerate self-consistent worlds	Clingo
<b>Well-Founded Semantics (WFS)</b>	three-valued: true / false / <b>undefined</b>	<code>WellFoundedSolver</code>

These correspond directly to profile classes shipped in `typedlogic.profiles` (`OpenWorld`, `ClosedWorld`, `WellFoundedSemantics`, `ClassicPrologNegationAsFailure`, `AnswerSetProgramming`), which we revisit at the end.

#### 3.1.1 Setup

This notebook uses the Z3, Clingo, and ProbLog integrations:

```
pip install 'typedlogic[z3,clingo,problog]'
```

Z3 is a classical first-order theorem prover — it embodies the **open-world, monotonic** view. Clingo is an **Answer Set Programming** solver — the **closed-world, non-monotonic** view. Putting them side by side is the whole point; `WellFoundedSolver` (built in) adds the three-valued middle ground.

Each theory lives in a `.tlog` file next to this notebook. Two tiny helpers let us *show* a theory (with syntax highlighting) and *load* it for a solver.

```
from pathlib import Path
from IPython.display import Code

from typedlogic import Term, Not
from typedlogic.parsers.tlog_parser import TLogParser
from typedlogic.integrations.solvers.z3 import Z3Solver
from typedlogic.integrations.solvers.clingo import ClingoSolver
from typedlogic.integrations.solvers.wellfounded import WellFoundedSolver, NegativeCycleError

HERE = Path("semantics") # the .tlog theory files live here

def show(name):
 "Display a .tlog theory with syntax highlighting."
 return Code(filename=str(HERE / name), language="prolog")

def load(name):
```

```
"Parse a .tlog theory file into a Theory."
return TLogParser().parse((HERE / name).read_text())
```

### 3.1.2 1. Open world vs closed world

Consider a tiny flight database. There is a direct flight SFO→JFK and JFK→LHR. There is **no** row for SFO→LHR.

```
show('flights.tlog')
```

```
pred DirectFlight(src: str, dst: str).

DirectFlight("sfo", "jfk").
DirectFlight("jfk", "lhr").
```

**The question:** *Is there a direct flight from SFO to LHR?* We never asserted one. What should a reasoner say?

#### Closed world (Clingo)

A closed-world reasoner treats the database as complete. Anything not derivable is taken to be false. Clingo returns exactly one model — the facts we gave it — and `DirectFlight(sfo, lhr)` is simply absent, i.e. **false**.

```
solver = ClingoSolver()
solver.add(load("flights.tlog"))

model = list(solver.models())[0]
present = {t.values for t in model.iter_retrieve("DirectFlight")}
print("Direct flights in the (single) closed-world model:")
for t in sorted(present):
 print(" ", t)
print()
print("Is DirectFlight(sfo, lhr) in the model?", ("sfo", "lhr") in present)
print("=> Under CWA, 'no row' means FALSE.")
```

```
Direct flights in the (single) closed-world model:
('jfk', 'lhr')
('sfo', 'jfk')

Is DirectFlight(sfo, lhr) in the model? False
=> Under CWA, 'no row' means FALSE.
```

#### Open world (Z3)

A classical, open-world reasoner treats the database as *partial*. The absence of SFO→LHR tells us nothing. We can prove this operationally: a fact is **entailed** only if its negation is *inconsistent* with the knowledge base. So we check satisfiability **both ways**.

```
flights = load("flights.tlog")

KB + DirectFlight(sfo,lhr): is it consistent?
s = Z3Solver(); s.add(flights); s.add(Term("DirectFlight", "sfo", "lhr"))
with_flight = s.check().satisfiable

KB + NOT DirectFlight(sfo,lhr): is it consistent?
s = Z3Solver(); s.add(flights); s.add(Not(Term("DirectFlight", "sfo", "lhr")))
without_flight = s.check().satisfiable

print("KB + DirectFlight(sfo,lhr) is satisfiable:", with_flight)
print("KB + ~DirectFlight(sfo,lhr) is satisfiable:", without_flight)
print()
print("Both are satisfiable => neither the flight nor its absence is entailed.")
print("=> Under OWA, 'no row' means UNKNOWN.")
```

```
KB + DirectFlight(sfo,lhr) is satisfiable: True
KB + ~DirectFlight(sfo,lhr) is satisfiable: True

Both are satisfiable => neither the flight nor its absence is entailed.
=> Under OWA, 'no row' means UNKNOWN.
```

This is the fundamental split:

- **Closed world:** what is not known to be true is **false**. Great for databases, configuration, policy checks — settings where you *do* have all the relevant facts.
- **Open world:** what is not known to be true is **unknown**. Right for the semantic web, ontologies, and incomplete knowledge — where new facts can always arrive (this is the OWL/RDF world; see the [OWL-DL tutorial](#)).

Neither is "correct" in the abstract; they answer different questions. The bugs happen when you assume one and your tool implements the other.

### 3.1.3 2. Negation as failure, and non-monotonicity

The closed-world assumption gives us a *cheap* form of negation: `not p` succeeds precisely when we **fail to prove** `p`. This is **negation as failure (NAF)** — the negation of Prolog, Datalog, and ASP. In TLog, `not` is negation-as-failure, while `~` is classical/strong negation.

The classic use is **default reasoning**: *birds fly, unless they are abnormal*.

```
show('birds.tlog')
```

```
pred Bird(name: str).
pred Penguin(name: str).
pred Abnormal(name: str).
pred Flies(name: str).

Bird("tweety").
Bird("opus").
Penguin("opus").

/// Penguins are abnormal birds.
Abnormal(x) :- Penguin(x).

/// A bird flies UNLESS it can be shown abnormal.
/// not here is negation-as-failure.
Flies(x) :- Bird(x), not Abnormal(x).
```

Let's ask Clingo who flies.

```
def who_flies(extra_facts=()):
 s = ClingoSolver()
 s.add(load("birds.tlog"))
 for f in extra_facts:
 s.add(f)
 model = list(s.models())[0]
 return sorted(t.values[0] for t in model.iter_retrieve("Flies"))

print("Who flies by default?", who_flies())
```

```
Who flies by default? ['tweety']
```

`tweety` flies; `opus` does not (it's a penguin, hence abnormal). Nobody told us `tweety` is *normal* — the reasoner **assumed** it because it couldn't prove otherwise. That assumption is the closed world at work.

Now the crucial property. Watch what happens when we *add* a fact.

```
print("Before: ", who_flies())
print("Add Penguin(tweety) ...")
print("After: ", who_flies(extra_facts=[Term("Penguin", "tweety")]))
print()
print("Learning MORE made us conclude LESS. This is NON-MONOTONICITY.")
```

```
Before: ['tweety']
Add Penguin(tweety) ...
After: []

Learning MORE made us conclude LESS. This is NON-MONOTONICITY.
```

Adding knowledge **retracted** a conclusion. Classical logic can never do this: it is **monotonic** — the set of entailed facts only ever grows. NAF buys expressiveness (defaults, exceptions, common-sense reasoning) at the price of monotonicity.

### The same rule, read classically (Z3)

What if we read "birds fly unless abnormal" with *classical* negation ( $\sim$ ) instead of NAF?

```
show('birds_classical.tlog')
```

```
pred Bird(name: str).
pred Abnormal(name: str).
pred Flies(name: str).

Bird("tweety").

/// The SAME rule, read with classical (strong) negation ~ instead of
/// negation-as-failure. Now ~Abnormal(x) must be
/// *proven*, not merely assumed.
Flies(x) :- Bird(x), ~Abnormal(x).
```

Under the open world, we *cannot* conclude `Flies(tweety)` from `Bird(tweety)` alone — because `~Abnormal(tweety)` is not entailed (it's merely unknown). We have to supply it.

```
bc = load("birds_classical.tlog")

Is Flies(tweety) entailed from Bird(tweety) alone?
s = Z3Solver(); s.add(bc); s.add(Not(Term("Flies", "tweety")))
print("KB + ~Flies(tweety) satisfiable?", s.check().satisfiable,
 " => Flies(tweety) NOT entailed (abnormality is unknown, not false)")

Now assert tweety is not abnormal, classically:
s = Z3Solver(); s.add(bc)
s.add(Not(Term("Abnormal", "tweety")))
s.add(Not(Term("Flies", "tweety")))
print("KB + ~Abnormal(tweety) + ~Flies(tweety) satisfiable?", s.check().satisfiable,
 " => now Flies(tweety) IS entailed")

KB + ~Flies(tweety) satisfiable? True => Flies(tweety) NOT entailed (abnormality is unknown, not false)
KB + ~Abnormal(tweety) + ~Flies(tweety) satisfiable? False => now Flies(tweety) IS entailed
```

So the *very same English sentence* means different things depending on the negation you pick:

- **NAF** (not Abnormal): "assume normal unless proven abnormal" → concludes `Flies(tweety)` immediately, but retracts it later. Non-monotonic.
- **Classical** ( $\sim$ Abnormal): "flies only if *provably* not abnormal" → stays agnostic until you supply the missing premise. Monotonic.

Choosing the wrong one is a real and common modelling bug.

### 3.1.4 3. When NAF loops: stable models (ASP)

Default reasoning was *stratified* — negation never depended on itself. Once we allow negation to loop back on itself, "what does `not p` mean" stops having an obvious answer, and the different semantics genuinely diverge. This is where **ASP** and **WFS** part ways.

**Answer Set Programming** answers with **stable models** (answer sets): each is a self-consistent world — a set of atoms that reproduces exactly itself when you apply the rules under its own negative assumptions. There can be several, or none.

### Even loop → two worlds

```
show('even_loop.tlog')
```

```
pred p().
pred q().

/// p holds unless q does, and q holds unless p does.
p() :- not q().
q() :- not p().
```

"p holds unless q does, and q holds unless p does." Two coherent stories: *p and not q*, or *q and not p*.

```
solver = ClingoSolver()
solver.add(load("even_loop.tlog"))
models = list(solver.models())
print("Number of stable models (answer sets):", len(models))
for i, m in enumerate(models):
 print(f" model {i}: ", sorted(str(t) for t in m.ground_terms))
```

```
Number of stable models (answer sets): 2
 model 0: ['q']
 model 1: ['p']
```

#### Odd loop → no world

```
show('odd_loop.tlog')
```

```
pred p().

/// p holds if p does not: a paradoxical negative self-loop.
p() :- not p().
```

"p holds if p does not." No set of atoms can reproduce itself: assume `p` false and the rule forces it true; assume it true and the rule no longer fires. **Zero** stable models — the program is *inconsistent* in ASP.

```
solver = ClingoSolver()
solver.add(load("odd_loop.tlog"))
print("Satisfiable?", solver.check().satisfiable)
print("Number of stable models:", len(list(solver.models())))
print("=> An odd negative loop has NO answer set.")
```

```
Satisfiable? False
Number of stable models: 0
=> An odd negative loop has NO answer set.
```

Stable-model semantics is decisive and credulous: it hands you every self-consistent world, or refuses when none exists. That's ideal for combinatorial search (planning, configuration, graph coloring), but the "all or nothing" reaction to a loop can be surprisingly brittle.

#### 3.1.5 4. A gentler answer: well-founded semantics

**Well-founded semantics (WFS)** refuses to guess. Instead of enumerating worlds it computes a *single three-valued* model: every atom is **true**, **false**, or **undefined**. Paradoxical loops land in *undefined* rather than exploding into many models or none.

typedlogic ships a `WellFoundedSolver` for exactly this. Its default `native` backend computes the well-founded model directly (via the standard Van Gelder alternating fixpoint) and reports the three-valued result. Run it on the same loops that gave ASP two models / zero models:

```
def well_founded(name):
 solver = WellFoundedSolver() # backend="native" by default
 solver.add(load(name))
 m = solver.model()
 true = sorted(str(t) for t in m.ground_terms)
 undefined = sorted(str(t) for t in m.undefined_terms)
 return true, undefined
```

```
for name in ["even_loop.tlog", "odd_loop.tlog"]:
 true, undefined = well_founded(name)
 print(f"{name:16} TRUE={true} UNDEFINED={undefined}")
```

```
even_loop.tlog TRUE=[] UNDEFINED=['p', 'q']
```

```
odd_loop.tlog TRUE=[] UNDEFINED=['p']
```

Compare the two semantics on the identical programs:

Program	ASP (stable models)	WFS (well-founded model)
<code>p:-not q. q:-not p.</code>	<b>two</b> models: <code>{p}</code> and <code>{q}</code>	one model: <code>p</code> , <code>q</code> both <b>undefined</b>
<code>p:-not p.</code>	<b>no</b> model (inconsistent)	one model: <code>p</code> <b>undefined</b>

ASP is *credulous and brave*: it commits to concrete worlds and enumerates them. WFS is *skeptical and cautious*: where the program is genuinely ambiguous or paradoxical, it declines to decide and marks the atom undefined — always yielding exactly one model, computable in polynomial time.

#### "Should we just wrap an existing engine?"

A fair question — WFS is well studied and has mature engines. The honest answer is *it depends on the fragment*, and the `WellFoundedSolver` backends make the trade-off concrete:

- The `problog` backend delegates to [ProbLog](#), a mature, wrapped engine. But ProbLog implements only the **two-valued** restriction of WFS: it agrees with the native backend on stratified programs, and **refuses** genuinely three-valued programs with a `NegativeCycle` error rather than reporting *undefined*.
- The `xsb` backend is meant to drive [XSB Prolog](#) — the reference SLG/tabling engine — for large or deeply-recursive programs (requires the external `xsb` binary). It is **experimental and not yet validated against a live XSB**, so prefer `native` for now.
- The `native` backend has no dependencies and returns the full three-valued model; it is the right tool for teaching and modestly-sized programs (grounding is naive, so it is not meant to compete with XSB on scale).

Watch ProbLog agree on the stratified default program, then refuse the negative loop:

```
On a stratified program, the wrapped ProbLog engine and the native engine agree exactly.
native = WellFoundedSolver(backend="native"); native.add(load("birds.tlog"))
problog = WellFoundedSolver(backend="problog"); problog.add(load("birds.tlog"))
native_flies = sorted(str(t) for t in native.model().iter_retrieve("Flies"))
problog_flies = sorted(str(t) for t in problog.model().iter_retrieve("Flies"))
print("native Flies:", native_flies)
print("problog Flies:", problog_flies)
print("agree?", native_flies == problog_flies)
print()
```

```
On the even loop, ProbLog refuses rather than reporting 'undefined'.
try:
 s = WellFoundedSolver(backend="problog"); s.add(load("even_loop.tlog"))
 s.model()
except NegativeCycleError as e:
 print("problog on even_loop -> NegativeCycleError:")
 print(" ", str(e).split(".")[0] + ".")
```

```
native Flies: ['Flies(tweety)']
problog Flies: ['Flies(tweety)']
agree? True
```

```
problog on even_loop -> NegativeCycleError:
The 'problog' backend only supports programs with a total (two-valued) well-founded model; this program is genuinely three-valued.
```

That is the whole reason the native three-valued backend exists: no pip-installable engine returns `undefined` for these programs. WFS and the other semantics agree whenever the program is well-behaved (stratified) — here is the `tweety/opus` default program under all three, for confirmation:

```
clingo (ASP), WFS native, WFS via problog -- all agree on the stratified program.
clingo = ClingoSolver(); clingo.add(load("birds.tlog"))
asp_flies = sorted(t.values[0] for t in list(clingo.models())[0].iter_retrieve("Flies"))

print("ASP (clingo) Flies:", asp_flies)
print("WFS (native) Flies:", native_flies)
print("WFS (problog) Flies:", problog_flies)
print()
print("All agree: on stratified programs WFS, ASP and Datalog coincide.")
```

```
ASP (clingo) Flies: ['tweety']
WFS (native) Flies: ['Flies(tweety)']
WFS (problog) Flies: ['Flies(tweety)']

All agree: on stratified programs WFS, ASP and Datalog coincide.
```

### 3.1.6 Putting it together: the typedlogic profiles

`typedlogic` reifies these semantic choices as **profile** classes in `typedlogic.profiles`, so a theory or solver can declare which assumptions it makes. The hierarchy encodes the relationships we just demonstrated — e.g. `WellFoundedSemantics` and `ClassicPrologNegationAsFailure` are both kinds of `ClosedWorld`, and `AnswerSetProgramming` is a `WellFoundedSemantics` that also permits `MultipleModelSemantics` (the several answer sets we saw). Each solver advertises its profile:

```
from typedlogic.profiles import (
 OpenWorld, ClosedWorld, SingleModelSemantics, MultipleModelSemantics,
)

solvers = [
 ("Z3 (classical FOL)", Z3Solver()),
 ("Clingo (ASP)", ClingoSolver()),
 ("WellFoundedSolver", WellFoundedSolver()),
]

hdr = f"{'solver':22} {'closed?':8} {'open?':6} {'single?':8} {'multi?':7}"
print(hdr); print("-" * len(hdr))
for label, s in solvers:
 p = s.profile
 print(f"{'label':22} {'str(p.impl(ClosedWorld))':8} {'str(p.impl(OpenWorld))':6} "
 f"{'str(p.impl(SingleModelSemantics))':8} {'str(p.impl(MultipleModelSemantics))':7}")
```

solver	closed?	open?	single?	multi?
Z3 (classical FOL)	False	True	False	True
Clingo (ASP)	True	False	False	True
WellFoundedSolver	True	False	True	False

#### Takeaways

- **Open vs closed world** decides what *silence* means: unknown, or false. Pick by whether your data is complete.
- **Negation as failure** gives cheap, powerful default reasoning — and with it **non-monotonicity**: new facts can retract old conclusions. Classical negation stays monotonic but can't express defaults.
- When negation loops, **ASP** enumerates stable worlds (two, or none), while **WFS** computes a single three-valued model that marks the ambiguity **undefined**. They coincide on stratified programs.
- The `WellFoundedSolver` gives you WFS directly, with a native three-valued engine, a wrapped Problog backend for the two-valued fragment, and an (experimental) XSB backend for scale.
- `typedlogic.profiles` lets you *name* the assumption you rely on, so a mismatch between what you meant and what your solver does becomes explicit rather than a silent bug.

#### Where to go next

- **Multiple models** — the [multiple-models notebook](#) explores enumerating worlds with Clingo in more depth.



- **Open-world modelling** — the [OWL-DL tutorial](#) is the open-world assumption applied to ontologies.
- **TLog syntax** — the [TLog parser docs](#) document the text syntax used for the theories here.

## 3.2 Multiple models

```
from typing import Union
from pydantic import BaseModel
```

```
class Organization(BaseModel):
 name: str
 address: str
```

```
class Person(BaseModel):
 name: str
 address: str
 year_of_birth: int
```

```
class Device(BaseModel):
 name: str
 manufactured_by: Union[Person, Organization]
```

```
my_device = Device(name="my_device", manufactured_by={"name": "John Doe", "address": "123 Main St.", "year_of_birth": 1980})
type(my_device.manufactured_by)
```

```
__main__.Person
```

```
my_device = Device(name="my_device", manufactured_by={"name": "John Doe", "address": "123 Main St."})
type(my_device.manufactured_by)
```

```
__main__.Organization
```

```
try:
 my_device.manufactured_by.year_of_birth = 1980
except ValueError as e:
 print(e)
```

```
"Organization" object has no field "year_of_birth"
```

```
my_device_dict = my_device.model_dump()
my_device_dict
```

```
{'name': 'my_device',
 'manufactured_by': {'name': 'John Doe', 'address': '123 Main St.'}}
```

```
import typedlogic.theories.jsonlog.loader as jsonlog_loader

facts = list(jsonlog_loader.generate_from_object(my_device_dict))
```

```
facts
```

```
[NodeIsObject(loc='/'),
 ObjectNodeLookup(loc='/', key='name', member='/name/'),
 NodeIsLiteral(loc='/name/'),
 NodeStringValue(loc='/name/', value='my_device'),
 ObjectNodeLookup(loc='/', key='manufactured_by', member='/manufactured_by/'),
 NodeIsObject(loc='/manufactured_by/'),
 ObjectNodeLookup(loc='/manufactured_by/', key='name', member='/manufactured_by/name/'),
 NodeIsLiteral(loc='/manufactured_by/name/'),
 NodeStringValue(loc='/manufactured_by/name/', value='John Doe'),
 ObjectNodeLookup(loc='/manufactured_by/', key='address', member='/manufactured_by/address/'),
 NodeIsLiteral(loc='/manufactured_by/address/'),
 NodeStringValue(loc='/manufactured_by/address/', value='123 Main St.')]]
```



## 4. Concepts

---

### 4.1 typed-logic: Bridging Formal Logic and Typed Python

---

typed-logic is a Python package that allows Python data models to be augmented using formal logical statements, which are then interpreted by *solvers* which can reason over combinations of programs and data allowing for *satisfiability checking*, and the generation of new data.

Currently the solvers supported are:

- [Z3](#)
- [Souffle](#)
- [Clingo](#)
- [Prover9](#)
- [Snakelog](#)
- [ProbLog](#)

With support for more solvers (Vampire, OWL reasoners, etc.) planned.

typed-logic is aimed primarily at software developers and data modelers who are logic-curious, but don't necessarily have a background in formal logic. It is especially aimed at Python developers who like to use lightweight ways of ensuring program and data correctness, such as Pydantic for data and mypy for type checking of programs.

#### 4.1.1 Key Features

---

- Write logical axioms and rules using familiar Python syntax
- Benefit from strong typing and mypy validation
- Seamless integration with logic programming engines
- Support for various solvers, including Z3 and Souffle
- Compatible with popular Python libraries like Pydantic

#### 4.1.2 Why TypedLogic?

---

TypedLogic combines the best of both worlds: the expressiveness and familiarity of Python with the power of formal logic and fast logic programming engines. This unique approach allows developers to:

1. Write more maintainable and less error-prone logical rules
2. Catch type-related errors early in the development process
3. Seamlessly integrate logical reasoning into existing Python projects
4. Leverage the performance of specialized logic engines without sacrificing the Python ecosystem

Get started with TypedLogic and experience a new way of combining logic programming with strongly typed Python!

## 4.2 Core Concepts

---

TypedLogic is built around several core concepts that blend logical programming with typed Python. Understanding these concepts is crucial for effectively using the library.

### 4.2.1 Use Python idioms to define your data structures

Facts are the basic units of information in TypedLogic. They are represented as Python classes that inherit from both `pydantic.BaseModel` and `FactMixin`. This approach allows you to define strongly-typed facts with automatic validation.

Example:

```
from pydantic import BaseModel
from typedlogic import FactMixin

PersonID = str
PetID = str

class Person(BaseModel, FactMixin):
 name: PersonID

class PersonAge(BaseModel, FactMixin):
 name: PersonID
 age: int

class Pet(BaseModel, FactMixin):
 name: PetID

class PetSpecies(BaseModel, FactMixin):
 name: PetID
 species: str

class OwnsPet(BaseModel, FactMixin):
 person: PersonID
 pet: PetID
```

Note in many logic frameworks these definitions don't have a direct translation, but in sorted logics, these may correspond to [predicate definitions](#).

### 4.2.2 Axioms

Axioms are logical rules or statements that define relationships between facts. In TypedLogic, axioms are defined using Python functions [decorated](#) with `@axiom`.

Example:

```
from typedlogic.decorators import axiom

@axiom
def constraints(person: PersonID, pet: PetID):
 if OwnsPet(person=person, pet=pet):
 assert Person(name=person) and Pet(name=pet)
```

Note that the Python code in the function definition is **never executed**. The code must be in a [subset of python](#) that is translated to logic statements.

You can also derive new facts from axioms:

```
class SameOwner(BaseModel, FactMixin):
 pet1: PetID
 pet2: PetID

from typedlogic.decorators import axiom

@axiom
def entail_same_owner(person: PersonID, pet1: PetID, pet2: PetID):
 if OwnsPet(person=person, pet=pet1) and OwnsPet(person=person, pet=pet2):
 assert SameOwner(pet1=pet1, pet2=pet2)
```

### 4.2.3 Generators

Generators like `gen1`, `gen2`, etc., are used within axioms to create typed placeholders for variables. They help maintain type safety while defining logical rules.

You can use generators in combination with Python `all` and `any` functions to express quantified sentences:

```
@axiom
def entail_same_owner():
 assert all(SameOwner(pet1=pet1, pet2=pet2)
```

```
for person, pet1, pet2 in gen3(PersonID, PetID, PetID)
if OwnsPet(person=person, pet=pet1) and OwnsPet(person=person, pet=pet2))
```

See [Generators](#) for more information.

## 4.2.4 Solvers

TypedLogic supports multiple [solvers](#), including Z3 and Souffle. Solvers are responsible for reasoning over the facts and axioms to derive new information or check for consistency.

Example (using the [Z3 solver](#)):

```
from typedlogic.integrations.solvers.z3 import Z3Solver

solver = Z3Solver()
solver.add(theory)
result = solver.check()
```

## 4.2.5 Theories

A [Theory](#) in TypedLogic is a collection of predicate definitions, facts, and axioms. It represents a complete knowledge base that can be reasoned over.

Example:

```
from typedlogic import Theory

theory = Theory(
 name="family_relationships",
 predicate_definitions=[...],
 sentence_groups=[...],
)
```

## 4.3 Data Model

Data model for the typed-logic framework.

### Overview

This module defines the core classes and structures used to represent logical constructs such as sentences, terms, predicates, and theories. It is based on the Common Logic Interchange Format (CLIF) and the Common Logic Standard (CL), with additions to make working with simple type systems easier.

Logical axioms are called [sentences](#) which organized into [theories](#)., which can be loaded into a [solver](#).

While one of the goals of typed-logic is to be able to write logic intuitively in Python, this data model is independent of the mapping from the Python language to the logic language; it can be used independently of the python syntax.

Here is an example:

```
>>> from typedlogic import Term, Forall, Implies
>>> x = Variable('x')
>>> y = Variable('y')
>>> pdef = PredicateDefinition(predicate='FriendOf',
... arguments={'x': 'str', 'y': 'str'}),
>>> theory = Theory(
... name="My theory",
... predicate_definitions=[pdef],
...)
>>> s = Forall([x, y],
... Implies(Term('friend_of', x, y),
... Term('friend_of', y, x)))
>>> theory.add(s)
```

### Classes

#### 4.3.1 PredicateDefinition `dataclass`

Defines the name and arguments of a predicate.

Example:

```
>>> pdef = PredicateDefinition(predicate='FriendOf',
... arguments={'x': 'str', 'y': 'str'})
```

The arguments are mappings between variable names and types. You can use either base types (e.g. 'str', 'int', 'float') or custom types.

Custom types should be defined in the theory's `type_definitions` attribute.

```
>>> pdef = PredicateDefinition(predicate='FriendOf',
... arguments={'x': 'Person', 'y': 'Person'})
>>> theory = Theory(
... name="My theory",
... type_definitions={'Person': 'str'},
... predicate_definitions=[pdef],
...)
```

Model:

```
classDiagram
class PredicateDefinition {
+String predicate
+Dict arguments
+String description
+Dict metadata
}
PredicateDefinition --> "*" PredicateDefinition : parents
```

### Source code in `src/typedlogic/datamodel.py`

```

49 @dataclass
50 class PredicateDefinition:
51 """
52 Defines the name and arguments of a predicate.
53
54 Example:
55
56 >>> pdef = PredicateDefinition(predicate='FriendOf',
57 ... arguments={'x': 'str', 'y': 'str'})
58
59 The arguments are mappings between variable names and types.
60 You can use either base types (e.g. 'str', 'int', 'float') or custom types.
61
62 Custom types should be defined in the theory's `type_definitions` attribute.
63
64 >>> pdef = PredicateDefinition(predicate='FriendOf',
65 ... arguments={'x': 'Person', 'y': 'Person'})
66 >>> theory = Theory(
67 ... name="My theory",
68 ... type_definitions={'Person': 'str'},
69 ... predicate_definitions=[pdef],
70 ...)
71
72 Model:
73
74 ``mermaid
75 classDiagram
76 class PredicateDefinition {
77 +String predicate
78 +Dict arguments
79 +String description
80 +Dict metadata
81 }
82 PredicateDefinition --> " " PredicateDefinition : parents
83 """
84
85 predicate: str
86 arguments: Dict[str, str]
87 description: Optional[str] = None
88 metadata: Optional[Dict[str, Any]] = None
89 parents: Optional[List[str]] = None
90 python_class: Optional[Type] = None
91
92 def argument_base_type(self, arg: str) -> str:
93 typ = self.arguments[arg]
94 try:
95 import pydantic
96
97 if isinstance(typ, pydantic.fields.FieldInfo):
98 typ = typ.annotation
99 except ImportError:
100 pass
101 return str(typ)
102
103 @classmethod
104 def from_class(cls, python_class: Type) -> "PredicateDefinition":
105 """
106 Create a predicate definition from a python class
107
108 :param predicate_class:
109 :return:
110 """
111 return PredicateDefinition(
112 predicate=python_class.__name__,
113 arguments={k: v for k, v in python_class.__annotations__.items()},
114)
115
116

```

`from_class(python_class)` classmethod

Create a predicate definition from a python class



**Parameters:**

Name	Type	Description	Default
<code>predicate_class</code>			<i>required</i>

**Returns:**

Type	Description
<code>PredicateDefinition</code>	

**Source code in `src/typedlogic/datamodel.py`**

```

105 @classmethod
106 def from_class(cls, python_class: Type) -> "PredicateDefinition":
107 """
108 Create a predicate definition from a python class
109
110 :param predicate_class:
111 :return:
112 """
113 return PredicateDefinition(
114 predicate=python_class.__name__,
115 arguments={k: v for k, v in python_class.__annotations__.items()},
116)

```

### 4.3.2 Variable dataclass

A variable in a logical sentence.

```

>>> x = Variable('x')
>>> y = Variable('y')
>>> s = Forall([x, y],
... Implies(Term('friend_of', x, y),
... Term('friend_of', y, x)))

```

Variables can have domains (types) specified:

```

>>> x = Variable('x', domain='str')
>>> y = Variable('y', domain='str')
>>> z = Variable('z', domain='int')
>>> xa = Variable('xa', domain='int')
>>> ya = Variable('ya', domain='int')
>>> s = Forall([x, y, z],
... Implies(And(Term('ParentOf', x, y),
... Term('Age', x, xa),
... Term('Age', y, ya)),
... Term('OlderThan', x, y)))

```

The domains should be either base types or defined types in the theory's `type_definitions` attribute.

### Source code in `src/typedlogic/datamodel.py`

```

119 @dataclass
120 class Variable:
121 """
122 A variable in a logical sentence.
123
124 >>> x = Variable('x')
125 >>> y = Variable('y')
126 >>> s = Forall([x, y],
127 ... Implies(Term('friend_of', x, y),
128 ... Term('friend_of', y, x)))
129
130 Variables can have domains (types) specified:
131
132 >>> x = Variable('x', domain='str')
133 >>> y = Variable('y', domain='str')
134 >>> z = Variable('z', domain='int')
135 >>> xa = Variable('xa', domain='int')
136 >>> ya = Variable('ya', domain='int')
137 >>> s = Forall([x, y, z],
138 ... Implies(And(Term('ParentOf', x, y),
139 ... Term('Age', x, xa),
140 ... Term('Age', y, ya)),
141 ... Term('OlderThan', x, y)))
142
143 The domains should be either base types or defined types in the theory's `type_definitions` attribute.
144 """
145
146 name: str
147 domain: Optional[str] = None
148 constraints: Optional[List[str]] = None
149
150 def __eq__(self, other):
151 return isinstance(other, Variable) and self.name == other.name
152
153 def __str__(self):
154 return "?" + self.name
155
156 def __hash__(self):
157 return hash(self.name)
158
159 def as_sexpr(self) -> SExpression:
160 sexpr = [type(self).__name__, self.name]
161 if self.domain:
162 return sexpr + [self.domain]
163 else:
164 return sexpr
165
166 @staticmethod
167 def create(names: str) -> tuple['Variable', ...]:
168 return tuple(Variable(name.strip()) for name in names.split())
169

```

### 4.3.3 Sentence

Bases: `ABC`

Base class for logical sentences.

Do not use this class directly; use one of the subclasses instead.

Model:

```

classDiagram
 Sentence <|-- Term
 Sentence <|-- BooleanSentence
 Sentence <|-- QuantifiedSentence
 Sentence <|-- Extension

```

 **Source code in** `src/typedlogic/datamodel.py` 

```

172 class Sentence(ABC):
173 """
174 Base class for logical sentences.
175
176 Do not use this class directly; use one of the subclasses instead.
177
178 Model:
179
180 ```mermaid
181 classDiagram
182 Sentence <|-- Term
183 Sentence <|-- BooleanSentence
184 Sentence <|-- QuantifiedSentence
185 Sentence <|-- Extension
186 """
187
188 """
189 def __init__(self):
190 self._annotations = {}
191
192 def __and__(self, other):
193 return And(self, other)
194
195 def __or__(self, other):
196 return Or(self, other)
197
198 def __invert__(self):
199 return Not(self)
200
201 def __sub__(self):
202 return NegationAsFailure(self)
203
204 def __rshift__(self, other):
205 return Implies(self, other)
206
207 def __lshift__(self, other):
208 return Implied(self, other)
209
210 def __xor__(self, other):
211 return Xor(self, other)
212
213 def iff(self, other):
214 return Iff(self, other)
215
216 def __lt__(self, other):
217 return Term(operator.lt.__name__, self, other)
218
219 def __le__(self, other):
220 return Term(operator.le.__name__, self, other)
221
222 def __gt__(self, other):
223 return Term(operator.gt.__name__, self, other)
224
225 def __ge__(self, other):
226 return Term(operator.ge.__name__, self, other)
227
228 def __add__(self, other):
229 return Term(operator.add.__name__, self, other)
230
231 def __mul__(self, other):
232 # TODO: avoid circular import
233 from typedlogic.extensions.probablilistic import Probability, That
234 return Probability(self, That(other))
235
236 @property
237 def annotations(self) -> Dict[str, Any]:
238 """
239 Annotations for the sentence.
240
241 Annotations are always logically silent, but can be used to store metadata or other information.
242
243 :return:
244 """
245 return self._annotations or {}
246
247 def add_annotation(self, key: str, value: Any):
248 """
249 Add an annotation to the sentence
250
251 :param key:
252 :param value:
253 :return:
254 """
255 if not self._annotations:
256 self._annotations = {}
257 self._annotations[key] = value
258
259 def as_sexpr(self) -> SExpression:
260 raise NotImplementedError(f"type = {type(self)} // {self}")
261
262 @property
263 def arguments(self) -> List[Any]:
264 raise NotImplementedError(f"type = {type(self)} // {self}")

```

257  
258  
259  
260  
261  
262  
263  
264  
265  
266

`annotations` property

Annotations for the sentence.

Annotations are always logically silent, but can be used to store metadata or other information.

**Returns:**

Type	Description
<code>Dict[str, Any]</code>	

`add_annotation(key, value)`

Add an annotation to the sentence

**Parameters:**

Name	Type	Description	Default
<code>key</code>	<code>str</code>		<i>required</i>
<code>value</code>	<code>Any</code>		<i>required</i>

**Returns:**

Type	Description

**Source code in** `src/typedlogic/datamodel.py` 

```
249 def add_annotation(self, key: str, value: Any):
250 """
251 Add an annotation to the sentence
252
253 :param key:
254 :param value:
255 :return:
256 """
257 if not self._annotations:
258 self._annotations = {}
259 self._annotations[key] = value
```

### 4.3.4 Term

Bases: [Sentence](#)

An atomic part of a sentence.

A ground term is a term with no variables:

```
>>> t = Term('FriendOf', 'Alice', 'Bob')
>>> t
FriendOf(Alice, Bob)
>>> t.values
('Alice', 'Bob')
```

```
>>> t.is_ground
True
```

Keyword argument based initialization is also supported:

```
>>> t = Term('FriendOf', dict(about='Alice', friend='Bob'))
>>> t.values
('Alice', 'Bob')
>>> t.positional
False
```

Mappings:

- Corresponds to AtomicSentence in Common Logic

 **Source code in** `src/typedlogic/datamodel.py` 

```

291 class Term(Sentence):
292 """
293 An atomic part of a sentence.
294
295 A ground term is a term with no variables:
296
297 >>> t = Term('FriendOf', 'Alice', 'Bob')
298 >>> t
299 FriendOf(Alice, Bob)
300 >>> t.values
301 ('Alice', 'Bob')
302 >>> t.is_ground
303 True
304
305 Keyword argument based initialization is also supported:
306
307 >>> t = Term('FriendOf', dict(about='Alice', friend='Bob'))
308 >>> t.values
309 ('Alice', 'Bob')
310 >>> t.positional
311 False
312
313 Mappings:
314 - Corresponds to AtomicSentence in Common Logic
315 """
316
317 def __init__(self, predicate: str, *args, **kwargs):
318 self.predicate = predicate
319 if not args:
320 self.positional = None
321 bindings = {}
322 elif len(args) == 1 and isinstance(args[0], dict):
323 bindings = args[0]
324 self.positional = False
325 else:
326 bindings = {f"arg{i}": arg for i, arg in enumerate(args)}
327 self.positional = True
328 self.bindings = bindings
329 self._annotations = kwargs
330
331 @property
332 def is_constant(self):
333 """
334 :return: True if the term is a constant (zero arguments)
335 """
336 return not self.bindings
337
338 @property
339 def is_ground(self):
340 """
341 :return: True if none of the arguments are variables
342 """
343 return not any(isinstance(v, Variable) for v in self.bindings.values())
344
345 @property
346 def values(self) -> Tuple[Any, ...]:
347 """
348 Representation of the arguments of the term as a fixed-position tuples
349 :return:
350 """
351 return tuple(self.bindings.values())
352
353 @property
354 def variables(self) -> List[Variable]:
355 """
356 :return: All of the arguments that are variables
357 """
358 return [v for v in self.bindings.values() if isinstance(v, Variable)]
359
360 @property
361 def variable_names(self) -> List[str]:
362 return [v.name for v in self.bindings.values() if isinstance(v, Variable)]
363
364 def make_keyword_indexed(self, keywords: List[str]):
365 """
366 Convert positional arguments to keyword arguments
367 """
368 if self.positional:
369 self.bindings = {k: v for k, v in zip(keywords, self.bindings.values(), strict=False)}
370 self.positional = False
371
372 def __repr__(self):
373 if not self.bindings:
374 return f"{self.predicate}"
375 elif self.positional:
376 return f'{self.predicate}({", ".join(f"{v}" for v in self.bindings.values())})'
377 else:
378 return f'{self.predicate}({", ".join(f"{k}={v}" for k, v in self.bindings.items())})'
379
380 def __eq__(self, other):
381 # return isinstance(other, Term) and self.predicate == other.predicate and self.bindings == other.bindings
382 return isinstance(other, Term) and self.predicate == other.predicate and self.values == other.values
383
384 def __hash__(self):
385 return hash((self.predicate, tuple(self.values)))
386
387 def as_sexpr(self) -> SExpression:
388 return [self.predicate] + [as_sexpr(v) for v in self.bindings.values()]

```